

X-Ray Diagnostics: Pediatric Pneumonia

This notebook will develop a model to classify pediatric chest x-rays with convolutional neural networks.

Importing Downloaded Data

Importing Necessary Python Libraries

```
In [1]: import os, shutil
```

1. Splitting Original Directories into Train/Test/Validation Directories
2. Loading Directory Paths & Contents into Variables

```
In [2]: # Loading Directory Paths into Variables
original_normal = 'ORIGINAL_DATA/NORMAL'
original_pneumonia = 'ORIGINAL_DATA/PNEUMONIA'

new_dir = 'data/'

train_folder = os.path.join(new_dir, 'train')
train_normal = os.path.join(train_folder, 'normal')
train_pneumonia = os.path.join(train_folder, 'pneumonia')

test_folder = os.path.join(new_dir, 'test')
test_normal = os.path.join(test_folder, 'normal')
test_pneumonia = os.path.join(test_folder, 'pneumonia')

val_folder = os.path.join(new_dir, 'validation')
val_normal = os.path.join(val_folder, 'normal')
val_pneumonia = os.path.join(val_folder, 'pneumonia')

# Creating Split Directories
os.mkdir(new_dir)

os.mkdir(test_folder)
os.mkdir(test_normal)
os.mkdir(test_pneumonia)

os.mkdir(train_folder)
os.mkdir(train_normal)
os.mkdir(train_pneumonia)

os.mkdir(val_folder)
os.mkdir(val_normal)
os.mkdir(val_pneumonia)
```

```
In [3]: # Exploring Raw Source Data

# Number of Images in NORMAL Directory
imgs_normal = [file for file in os.listdir(
    original_normal) if file.endswith('.jpeg')]
print(len(imgs_normal), 'images in NORMAL directory')

# Number of Images in PNEUMONIA Directory
imgs_pneumonia = [file for file in os.listdir(
```

```
original_pneumonia) if file.endswith('.jpeg')]]  
print(len(imgs_pneumonia), 'images in PNEUMONIA directory')
```

1583 images in NORMAL directory
4273 images in PNEUMONIA directory

```
In [4]: # Copying Raw Data into Split Directories  
  
# train normal  
imgs = imgs_normal[:1200]  
for img in imgs:  
    origin = os.path.join(original_normal, img)  
    destination = os.path.join(train_normal, img)  
    shutil.copyfile(origin, destination)  
  
# test normal  
imgs = imgs_normal[1200:1383]  
for img in imgs:  
    origin = os.path.join(original_normal, img)  
    destination = os.path.join(test_normal, img)  
    shutil.copyfile(origin, destination)  
  
# validation normal  
imgs = imgs_normal[1383:]  
for img in imgs:  
    origin = os.path.join(original_normal, img)  
    destination = os.path.join(val_normal, img)  
    shutil.copyfile(origin, destination)  
  
# train pneumonia  
imgs = imgs_pneumonia[:3900]  
for img in imgs:  
    origin = os.path.join(original_pneumonia, img)  
    destination = os.path.join(train_pneumonia, img)  
    shutil.copyfile(origin, destination)  
  
# test pneumonia  
imgs = imgs_pneumonia[3900:4073]  
for img in imgs:  
    origin = os.path.join(original_pneumonia, img)  
    destination = os.path.join(test_pneumonia, img)  
    shutil.copyfile(origin, destination)  
  
# validation pneumonia  
imgs = imgs_pneumonia[4073:]  
for img in imgs:  
    origin = os.path.join(original_pneumonia, img)  
    destination = os.path.join(val_pneumonia, img)  
    shutil.copyfile(origin, destination)
```

```
In [5]: ## CELL INTENDED TO RE-ESTABLISH VARIABLES ##  
## FROM DEAD/RESTARTED KERNELS ##  
  
# Loading Directory Paths into Variables  
train_folder = 'data/train'  
train_normal = 'data/train/normal'  
train_pneumonia = 'data/train/pneumonia'  
  
test_folder = 'data/test'  
test_normal = 'data/test/normal'  
test_pneumonia = 'data/test/pneumonia'  
  
val_folder = 'data/validation'  
val_normal = 'data/validation/normal'  
val_pneumonia = 'data/validation/pneumonia'
```

```
In [6]: # Number of Images in Each Directory

a1 = len(os.listdir(train_normal))
a2 = len(os.listdir(train_pneumonia))
a = a1 + a2
b1 = len(os.listdir(test_normal))
b2 = len(os.listdir(test_pneumonia))
b = b1 + b2
c1 = len(os.listdir(val_normal))
c2 = len(os.listdir(val_pneumonia))
c = c1 + c2

print(a, 'images in train directory')
print(b, 'images in test directory')
print(c, 'images in validation directory')
```

5100 images in train directory
 356 images in test directory
 400 images in validation directory

Preprocessing Data

Importing Necessary Python Libraries

```
In [7]: import scipy
import numpy as np
from PIL import Image
from scipy import ndimage
from keras.preprocessing.image import (
    ImageDataGenerator, array_to_img,
    img_to_array, load_img)

np.random.seed(123)
```

Preprocess Part A

```
In [8]: # flow_from_directory Variables
targetimagesize_ = (150, 150)
trainbatchsize_ = a
testbatchsize_ = b
valbatchsize_ = c

# Reshape Data in train Directory
train_generator = ImageDataGenerator(
    rescale=1./255).flow_from_directory(
    train_folder,
    target_size = targetimagesize_,
    batch_size = trainbatchsize_)

# Reshape Data in test Directory
test_generator = ImageDataGenerator(
    rescale=1./255).flow_from_directory(
    test_folder,
    target_size = targetimagesize_,
    batch_size = testbatchsize_)

# Reshape Data in validation Directory
val_generator = ImageDataGenerator(
    rescale=1./255).flow_from_directory(
    val_folder,
    target_size = targetimagesize_,
    batch_size = valbatchsize_)
```

Found 5100 images belonging to 2 classes.
Found 356 images belonging to 2 classes.
Found 400 images belonging to 2 classes.

```
In [9]: # Load Dataset into Variables
train_images, train_labels = next(train_generator)
test_images, test_labels = next(test_generator)
val_images, val_labels = next(val_generator)
```

```
In [10]: # Exploring Final Datasets
m_train = train_images.shape[0]
m_test = test_images.shape[0]
m_val = val_images.shape[0]

print ("Number of training samples: " + str(m_train))
print ("Number of testing samples: " + str(m_test))
print ("Number of validation samples: " + str(m_val))
print ("train_images shape: " + str(train_images.shape))
print ("train_labels shape: " + str(train_labels.shape))
print ("test_images shape: " + str(test_images.shape))
print ("test_labels shape: " + str(test_labels.shape))
print ("val_images shape: " + str(val_images.shape))
print ("val_labels shape: " + str(val_labels.shape))
```

Number of training samples: 5100
Number of testing samples: 356
Number of validation samples: 400
train_images shape: (5100, 150, 150, 3)
train_labels shape: (5100, 2)
test_images shape: (356, 150, 150, 3)
test_labels shape: (356, 2)
val_images shape: (400, 150, 150, 3)
val_labels shape: (400, 2)

Preprocess Part B

```
In [11]: # Reshaping into 2-D Array

train_img = train_images.reshape(train_images.shape[0], -1)
test_img = test_images.reshape(test_images.shape[0], -1)
val_img = val_images.reshape(val_images.shape[0], -1)

print(train_img.shape)
print(test_img.shape)
print(val_img.shape)
```

(5100, 67500)
(356, 67500)
(400, 67500)

```
In [12]: # Loading y Variables as 2-D Array

train_y = np.reshape(train_labels[:,0], (trainbatchsize_,1))
test_y = np.reshape(test_labels[:,0], (testbatchsize_,1))
val_y = np.reshape(val_labels[:,0], (valbatchsize_,1))

input_shape_ = train_img.shape[1]
```

Modeling

Importing Necessary Python Libraries

```
In [13]: # Importing Python Libraries to Fit Models
from keras.models import Sequential
from keras.layers import (
```

```

Conv2D, Dense, Flatten, MaxPooling2D)
from keras import optimizers

# Importing Python Libraries for Analysis
import datetime
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import (
    classification_report, roc_curve, auc,
    confusion_matrix)

%matplotlib inline
np.random.seed(123)

```

Baseline Model

```

In [14]: # Building the Model

def Build_Baseline():
    model = Sequential()
    model.add(Dense(20, activation='relu',
                    input_shape=(input_shape_,)))
    model.add(Dense(7, activation='relu'))
    model.add(Dense(5, activation='relu'))
    model.add(Dense(1, activation='sigmoid'))

    # Compile model
    model.compile(optimizer='sgd',
                  loss='binary_crossentropy',
                  metrics=['acc'])

    return model

```

```

In [15]: # Timer Start
start = datetime.datetime.now()

```

```

In [16]: base = Build_Baseline()
basehist = base.fit(train_img,
                    train_y,
                    epochs=50,
                    batch_size=32,
                    validation_data=(val_img, val_y))

```

```

Epoch 1/50
160/160 [=====] - 2s 11ms/step - loss: 0.4558 - acc: 0.7996 - val_loss:
0.5014 - val_acc: 0.7825
Epoch 2/50
160/160 [=====] - 1s 9ms/step - loss: 0.3256 - acc: 0.8669 - val_loss: 1.
3702 - val_acc: 0.5200
Epoch 3/50
160/160 [=====] - 1s 8ms/step - loss: 0.2586 - acc: 0.8941 - val_loss: 0.
2888 - val_acc: 0.8700
Epoch 4/50
160/160 [=====] - 1s 8ms/step - loss: 0.2307 - acc: 0.9039 - val_loss: 0.
8482 - val_acc: 0.6125
Epoch 5/50
160/160 [=====] - 1s 8ms/step - loss: 0.1877 - acc: 0.9286 - val_loss: 0.
4494 - val_acc: 0.8100
Epoch 6/50
160/160 [=====] - 1s 8ms/step - loss: 0.1951 - acc: 0.9224 - val_loss: 0.
5184 - val_acc: 0.7875
Epoch 7/50
160/160 [=====] - 1s 8ms/step - loss: 0.1822 - acc: 0.9273 - val_loss: 0.
5070 - val_acc: 0.8000
Epoch 8/50
160/160 [=====] - 1s 8ms/step - loss: 0.1746 - acc: 0.9335 - val_loss: 0.
2193 - val_acc: 0.9025

```

Epoch 9/50
160/160 [=====] - 1s 8ms/step - loss: 0.1794 - acc: 0.9280 - val_loss: 0.4827 - val_acc: 0.7850
Epoch 10/50
160/160 [=====] - 1s 8ms/step - loss: 0.1758 - acc: 0.9329 - val_loss: 0.3105 - val_acc: 0.8750
Epoch 11/50
160/160 [=====] - 1s 8ms/step - loss: 0.1617 - acc: 0.9375 - val_loss: 0.4740 - val_acc: 0.7925
Epoch 12/50
160/160 [=====] - 1s 8ms/step - loss: 0.1738 - acc: 0.9357 - val_loss: 0.3565 - val_acc: 0.8500
Epoch 13/50
160/160 [=====] - 1s 8ms/step - loss: 0.1520 - acc: 0.9427 - val_loss: 0.4667 - val_acc: 0.8000
Epoch 14/50
160/160 [=====] - 1s 8ms/step - loss: 0.1501 - acc: 0.9406 - val_loss: 0.3021 - val_acc: 0.8775
Epoch 15/50
160/160 [=====] - 1s 8ms/step - loss: 0.1422 - acc: 0.9451 - val_loss: 0.3877 - val_acc: 0.8475
Epoch 16/50
160/160 [=====] - 1s 8ms/step - loss: 0.1423 - acc: 0.9455 - val_loss: 0.2056 - val_acc: 0.9125
Epoch 17/50
160/160 [=====] - 1s 8ms/step - loss: 0.1525 - acc: 0.9396 - val_loss: 0.2109 - val_acc: 0.9075
Epoch 18/50
160/160 [=====] - 1s 8ms/step - loss: 0.1332 - acc: 0.9520 - val_loss: 0.4235 - val_acc: 0.8200
Epoch 19/50
160/160 [=====] - 1s 8ms/step - loss: 0.1444 - acc: 0.9469 - val_loss: 0.2334 - val_acc: 0.9075
Epoch 20/50
160/160 [=====] - 1s 8ms/step - loss: 0.1354 - acc: 0.9469 - val_loss: 0.2266 - val_acc: 0.9050
Epoch 21/50
160/160 [=====] - 1s 8ms/step - loss: 0.1335 - acc: 0.9478 - val_loss: 0.2035 - val_acc: 0.9300
Epoch 22/50
160/160 [=====] - 1s 8ms/step - loss: 0.1369 - acc: 0.9480 - val_loss: 0.2167 - val_acc: 0.9150
Epoch 23/50
160/160 [=====] - 1s 8ms/step - loss: 0.1274 - acc: 0.9502 - val_loss: 0.2256 - val_acc: 0.9125
Epoch 24/50
160/160 [=====] - 1s 8ms/step - loss: 0.1318 - acc: 0.9492 - val_loss: 0.2837 - val_acc: 0.8850
Epoch 25/50
160/160 [=====] - 1s 8ms/step - loss: 0.1326 - acc: 0.9500 - val_loss: 0.2191 - val_acc: 0.9025
Epoch 26/50
160/160 [=====] - 1s 9ms/step - loss: 0.1297 - acc: 0.9504 - val_loss: 0.2172 - val_acc: 0.9125
Epoch 27/50
160/160 [=====] - 1s 8ms/step - loss: 0.1255 - acc: 0.9514 - val_loss: 0.2213 - val_acc: 0.9125
Epoch 28/50
160/160 [=====] - 1s 8ms/step - loss: 0.1239 - acc: 0.9551 - val_loss: 0.2329 - val_acc: 0.8950
Epoch 29/50
160/160 [=====] - 1s 8ms/step - loss: 0.1253 - acc: 0.9522 - val_loss: 0.2848 - val_acc: 0.8900
Epoch 30/50
160/160 [=====] - 1s 8ms/step - loss: 0.1234 - acc: 0.9522 - val_loss: 0.2428 - val_acc: 0.9000
Epoch 31/50
160/160 [=====] - 1s 8ms/step - loss: 0.1237 - acc: 0.9529 - val_loss: 0.5362 - val_acc: 0.7850
Epoch 32/50

```

160/160 [=====] - 1s 8ms/step - loss: 0.1182 - acc: 0.9576 - val_loss: 0.
3199 - val_acc: 0.8875
Epoch 33/50
160/160 [=====] - 1s 8ms/step - loss: 0.1201 - acc: 0.9563 - val_loss: 0.
1935 - val_acc: 0.9325
Epoch 34/50
160/160 [=====] - 1s 8ms/step - loss: 0.1194 - acc: 0.9537 - val_loss: 0.
3693 - val_acc: 0.8375
Epoch 35/50
160/160 [=====] - 1s 8ms/step - loss: 0.1149 - acc: 0.9573 - val_loss: 0.
2020 - val_acc: 0.9200
Epoch 36/50
160/160 [=====] - 1s 8ms/step - loss: 0.1112 - acc: 0.9588 - val_loss: 0.
2087 - val_acc: 0.9200
Epoch 37/50
160/160 [=====] - 1s 8ms/step - loss: 0.1088 - acc: 0.9598 - val_loss: 0.
1816 - val_acc: 0.9300
Epoch 38/50
160/160 [=====] - 1s 8ms/step - loss: 0.1086 - acc: 0.9582 - val_loss: 0.
7579 - val_acc: 0.7475
Epoch 39/50
160/160 [=====] - 1s 8ms/step - loss: 0.1131 - acc: 0.9582 - val_loss: 0.
4557 - val_acc: 0.8125
Epoch 40/50
160/160 [=====] - 1s 8ms/step - loss: 0.1145 - acc: 0.9580 - val_loss: 0.
1809 - val_acc: 0.9325
Epoch 41/50
160/160 [=====] - 1s 8ms/step - loss: 0.1062 - acc: 0.9625 - val_loss: 0.
4450 - val_acc: 0.8275
Epoch 42/50
160/160 [=====] - 1s 8ms/step - loss: 0.1147 - acc: 0.9571 - val_loss: 0.
2581 - val_acc: 0.8925
Epoch 43/50
160/160 [=====] - 1s 8ms/step - loss: 0.1114 - acc: 0.9600 - val_loss: 0.
2369 - val_acc: 0.9125
Epoch 44/50
160/160 [=====] - 1s 8ms/step - loss: 0.1061 - acc: 0.9588 - val_loss: 0.
1993 - val_acc: 0.9200
Epoch 45/50
160/160 [=====] - 1s 8ms/step - loss: 0.1086 - acc: 0.9590 - val_loss: 0.
2287 - val_acc: 0.9025
Epoch 46/50
160/160 [=====] - 1s 8ms/step - loss: 0.1016 - acc: 0.9612 - val_loss: 0.
2337 - val_acc: 0.9025
Epoch 47/50
160/160 [=====] - 1s 8ms/step - loss: 0.1011 - acc: 0.9653 - val_loss: 0.
1786 - val_acc: 0.9350
Epoch 48/50
160/160 [=====] - 1s 8ms/step - loss: 0.1037 - acc: 0.9610 - val_loss: 0.
3447 - val_acc: 0.8575
Epoch 49/50
160/160 [=====] - 1s 8ms/step - loss: 0.0976 - acc: 0.9643 - val_loss: 0.
1913 - val_acc: 0.9275
Epoch 50/50
160/160 [=====] - 1s 8ms/step - loss: 0.1044 - acc: 0.9610 - val_loss: 0.
2143 - val_acc: 0.9100

```

```

In [17]: # Timer End
end = datetime.datetime.now()
elapsed = end - start
print('Training Elapsed Time: {}'.format(elapsed))

```

Training Elapsed Time: 0:01:09.876972

Baseline Model Analysis

```

In [124... # Loading Variables for Analysis

results_train = base.evaluate(train_img, train_y)

```

```
results_test = base.evaluate(test_img, test_y)
```

```
pred_y = base.predict(test_img).ravel()
```

```
fpr_, tpr_, thresholds_ = roc_curve(test_y, pred_y)  
auc_ = auc(fpr_, tpr_)
```

```
160/160 [=====] - 1s 4ms/step - loss: 0.1165 - acc: 0.9537  
12/12 [=====] - 0s 2ms/step - loss: 0.2521 - acc: 0.9045
```



In [125...

```
# Evaluation Results  
print ('Train Results:', results_train)  
print ('Test Results:', results_test)
```

```
Train Results: [0.11654186993837357, 0.9537255167961121]  
Test Results: [0.25213193893432617, 0.9044944047927856]
```

In [126...

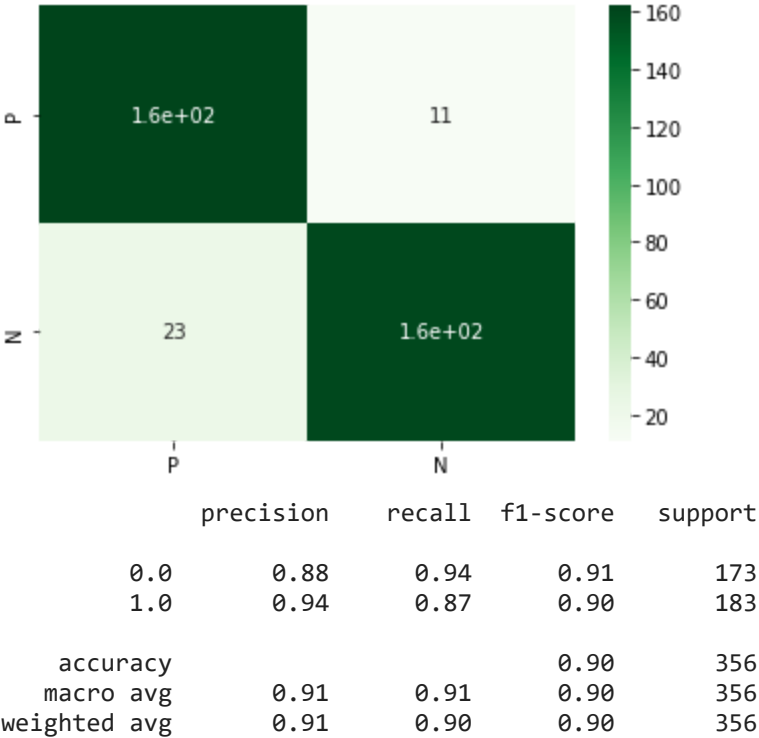
```
acc = basehist.history['acc']  
val_acc = basehist.history['val_acc']  
loss = basehist.history['loss']  
val_loss = basehist.history['val_loss']  
epochs = range(len(acc))  
plt.plot(epochs, acc, 'bo', label='Training Acc')  
plt.plot(epochs, val_acc, 'b', label='Validation Acc')  
plt.title('Training and Validation Accuracy')  
plt.legend()  
plt.figure()  
plt.plot(epochs, loss, 'bo', label='Training Loss')  
plt.plot(epochs, val_loss, 'b', label='Validation Loss')  
plt.title('Training and Validation Loss')  
plt.legend()  
plt.show()
```




```
In [127... pred_y = (base.predict(test_img).ravel() > 0.5).astype(int)
```

```
In [128... # Confusion Matrix
cm = confusion_matrix(test_y, pred_y)
f = sns.heatmap(cm, annot=True, cmap='Greens',
                xticklabels='PN', yticklabels='PN')
plt.show()


# Classification Report
print(classification_report(test_y, pred_y))
```



```
In [129... print('False Normal Rate:', (11/b2)*100)

False Normal Rate: 6.358381502890173
```

```
In [130... plt.figure(1)
plt.plot([0, 1], [0, 1], 'k--')
plt.plot(fpr_, tpr_,
        label='model(area = {:.3f})'.format(auc_))
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc='best')
plt.show()
```




In [131... auc_

Out[131... 0.9736567800625414

In [134... pred_y = (base.predict(test_img).ravel() > 0.95).astype(int)

In [135...
Confusion Matrix
cm = confusion_matrix(test_y, pred_y)
f = sns.heatmap(cm, annot=True, cmap='cubehelix',
xticklabels='PN', yticklabels='PN')
plt.show()

Classification Report
print(classification_report(test_y, pred_y))



In [136... print('False Normal Rate:', (2/b2)*100)

False Normal Rate: 1.1560693641618496

In [59]: base.save('XRAY_Baseline_Model.h5')

CNN Model

```
In [26]: # Building the Model

def Build_CNN():
    model = Sequential()

    model.add(Conv2D(32, (3, 3), activation='relu',
                     input_shape=(150, 150, 3)))
    model.add(MaxPooling2D((2, 2)))
    model.add(Conv2D(32, (4, 4), activation='relu'))
    model.add(MaxPooling2D((2, 2)))
    model.add(Conv2D(64, (3, 3), activation='relu'))
    model.add(MaxPooling2D((2, 2)))

    model.add(Flatten())
    model.add(Dense(64, activation='relu'))
    model.add(Dense(1, activation='sigmoid'))

    # Compile model
    model.compile(optimizer='sgd',
                  loss='binary_crossentropy',
                  metrics=['acc'])

    return model
```

```
In [27]: # Timer Start
start = datetime.datetime.now()
```

```
In [29]: cnn = Build_CNN()
history = cnn.fit(train_images,
                  train_y,
                  epochs=50,
                  batch_size=32,
                  validation_data=(val_images, val_y))
```

```
Epoch 1/50
160/160 [=====] - 73s 458ms/step - loss: 0.5162 - acc: 0.7727 - val_loss:
0.5916 - val_acc: 0.6425
Epoch 2/50
160/160 [=====] - 74s 460ms/step - loss: 0.3747 - acc: 0.8406 - val_loss:
0.3607 - val_acc: 0.8500
Epoch 3/50
160/160 [=====] - 74s 461ms/step - loss: 0.2678 - acc: 0.8957 - val_loss:
0.3012 - val_acc: 0.8650
Epoch 4/50
160/160 [=====] - 74s 463ms/step - loss: 0.2082 - acc: 0.9196 - val_loss:
0.2674 - val_acc: 0.8825
Epoch 5/50
160/160 [=====] - 74s 461ms/step - loss: 0.1899 - acc: 0.9259 - val_loss:
0.7132 - val_acc: 0.6750
Epoch 6/50
160/160 [=====] - 75s 470ms/step - loss: 0.1676 - acc: 0.9357 - val_loss:
0.2129 - val_acc: 0.9150
Epoch 7/50
160/160 [=====] - 74s 460ms/step - loss: 0.1559 - acc: 0.9443 - val_loss:
0.1949 - val_acc: 0.9175
Epoch 8/50
160/160 [=====] - 74s 461ms/step - loss: 0.1457 - acc: 0.9461 - val_loss:
0.3333 - val_acc: 0.8625
Epoch 9/50
160/160 [=====] - 72s 453ms/step - loss: 0.1410 - acc: 0.9484 - val_loss:
0.1715 - val_acc: 0.9325
Epoch 10/50
160/160 [=====] - 80s 497ms/step - loss: 0.1327 - acc: 0.9522 - val_loss:
0.1695 - val_acc: 0.9450
Epoch 11/50
160/160 [=====] - 74s 461ms/step - loss: 0.1261 - acc: 0.9531 - val_loss:
```

0.1835 - val_acc: 0.9250
Epoch 12/50
160/160 [=====] - 74s 463ms/step - loss: 0.1187 - acc: 0.9567 - val_loss:
0.2253 - val_acc: 0.9175
Epoch 13/50
160/160 [=====] - 76s 474ms/step - loss: 0.1135 - acc: 0.9584 - val_loss:
0.1566 - val_acc: 0.9375
Epoch 14/50
160/160 [=====] - 74s 461ms/step - loss: 0.1089 - acc: 0.9608 - val_loss:
0.2095 - val_acc: 0.9175
Epoch 15/50
160/160 [=====] - 71s 446ms/step - loss: 0.1048 - acc: 0.9602 - val_loss:
0.1643 - val_acc: 0.9450
Epoch 16/50
160/160 [=====] - 65s 408ms/step - loss: 0.0979 - acc: 0.9629 - val_loss:
0.2025 - val_acc: 0.9225
Epoch 17/50
160/160 [=====] - 65s 405ms/step - loss: 0.0998 - acc: 0.9641 - val_loss:
0.1763 - val_acc: 0.9425
Epoch 18/50
160/160 [=====] - 65s 407ms/step - loss: 0.0932 - acc: 0.9651 - val_loss:
0.1336 - val_acc: 0.9525
Epoch 19/50
160/160 [=====] - 65s 405ms/step - loss: 0.0901 - acc: 0.9667 - val_loss:
0.1328 - val_acc: 0.9625
Epoch 20/50
160/160 [=====] - 65s 406ms/step - loss: 0.0888 - acc: 0.9667 - val_loss:
0.1662 - val_acc: 0.9400
Epoch 21/50
160/160 [=====] - 65s 404ms/step - loss: 0.0851 - acc: 0.9676 - val_loss:
0.1287 - val_acc: 0.9600
Epoch 22/50
160/160 [=====] - 65s 405ms/step - loss: 0.0830 - acc: 0.9682 - val_loss:
0.1561 - val_acc: 0.9550
Epoch 23/50
160/160 [=====] - 65s 405ms/step - loss: 0.0755 - acc: 0.9716 - val_loss:
0.1499 - val_acc: 0.9475
Epoch 24/50
160/160 [=====] - 65s 404ms/step - loss: 0.0771 - acc: 0.9727 - val_loss:
0.1307 - val_acc: 0.9625
Epoch 25/50
160/160 [=====] - 65s 405ms/step - loss: 0.0729 - acc: 0.9731 - val_loss:
0.1397 - val_acc: 0.9550
Epoch 26/50
160/160 [=====] - 65s 407ms/step - loss: 0.0702 - acc: 0.9741 - val_loss:
0.1732 - val_acc: 0.9350
Epoch 27/50
160/160 [=====] - 65s 404ms/step - loss: 0.0676 - acc: 0.9751 - val_loss:
0.1148 - val_acc: 0.9625
Epoch 28/50
160/160 [=====] - 65s 408ms/step - loss: 0.0647 - acc: 0.9755 - val_loss:
0.2118 - val_acc: 0.9200
Epoch 29/50
160/160 [=====] - 65s 405ms/step - loss: 0.0646 - acc: 0.9765 - val_loss:
0.1289 - val_acc: 0.9550
Epoch 30/50
160/160 [=====] - 65s 405ms/step - loss: 0.0633 - acc: 0.9739 - val_loss:
0.1178 - val_acc: 0.9700
Epoch 31/50
160/160 [=====] - 65s 406ms/step - loss: 0.0566 - acc: 0.9796 - val_loss:
0.1282 - val_acc: 0.9575
Epoch 32/50
160/160 [=====] - 66s 410ms/step - loss: 0.0558 - acc: 0.9806 - val_loss:
0.1449 - val_acc: 0.9600
Epoch 33/50
160/160 [=====] - 65s 405ms/step - loss: 0.0520 - acc: 0.9814 - val_loss:
0.1240 - val_acc: 0.9650
Epoch 34/50
160/160 [=====] - 65s 406ms/step - loss: 0.0531 - acc: 0.9802 - val_loss:
0.3962 - val_acc: 0.8800

```

Epoch 35/50
160/160 [=====] - 72s 451ms/step - loss: 0.0531 - acc: 0.9820 - val_loss:
0.1303 - val_acc: 0.9575
Epoch 36/50
160/160 [=====] - 1269s 8s/step - loss: 0.0469 - acc: 0.9837 - val_loss:
0.1437 - val_acc: 0.9550
Epoch 37/50
160/160 [=====] - 60s 375ms/step - loss: 0.0503 - acc: 0.9818 - val_loss:
0.1247 - val_acc: 0.9550
Epoch 38/50
160/160 [=====] - 59s 366ms/step - loss: 0.0441 - acc: 0.9837 - val_loss:
0.1289 - val_acc: 0.9550
Epoch 39/50
160/160 [=====] - 58s 364ms/step - loss: 0.0469 - acc: 0.9827 - val_loss:
0.1221 - val_acc: 0.9575
Epoch 40/50
160/160 [=====] - 59s 369ms/step - loss: 0.0425 - acc: 0.9863 - val_loss:
0.2873 - val_acc: 0.9000
Epoch 41/50
160/160 [=====] - 62s 385ms/step - loss: 0.0404 - acc: 0.9857 - val_loss:
0.1403 - val_acc: 0.9475
Epoch 42/50
160/160 [=====] - 61s 383ms/step - loss: 0.0398 - acc: 0.9857 - val_loss:
0.1375 - val_acc: 0.9575
Epoch 43/50
160/160 [=====] - 59s 372ms/step - loss: 0.0384 - acc: 0.9855 - val_loss:
0.1390 - val_acc: 0.9550
Epoch 44/50
160/160 [=====] - 63s 392ms/step - loss: 0.0353 - acc: 0.9884 - val_loss:
0.1328 - val_acc: 0.9625
Epoch 45/50
160/160 [=====] - 66s 413ms/step - loss: 0.0392 - acc: 0.9863 - val_loss:
0.1209 - val_acc: 0.9600
Epoch 46/50
160/160 [=====] - 65s 406ms/step - loss: 0.0323 - acc: 0.9900 - val_loss:
0.1782 - val_acc: 0.9500
Epoch 47/50
160/160 [=====] - 66s 410ms/step - loss: 0.0421 - acc: 0.9849 - val_loss:
0.2284 - val_acc: 0.9350
Epoch 48/50
160/160 [=====] - 67s 417ms/step - loss: 0.0304 - acc: 0.9888 - val_loss:
0.1670 - val_acc: 0.9550
Epoch 49/50
160/160 [=====] - 66s 413ms/step - loss: 0.0295 - acc: 0.9890 - val_loss:
0.1363 - val_acc: 0.9650
Epoch 50/50
160/160 [=====] - 66s 413ms/step - loss: 0.0254 - acc: 0.9910 - val_loss:
0.1306 - val_acc: 0.9575

```

```

In [30]: # Timer End
end = datetime.datetime.now()
elapsed = end - start
print('Training Elapsed Time: {}'.format(elapsed))

```

Training Elapsed Time: 1:26:36.105971

CNN Model Analysis

```

In [137... # Loading Variables for Analysis

results_train = cnn.evaluate(train_images, train_y)
results_test = cnn.evaluate(test_images, test_y)

pred_y = cnn.predict(test_images).ravel()

fpr_, tpr_, thresholds_ = roc_curve(test_y, pred_y)
auc_ = auc(fpr_, tpr_)

```

```

160/160 [=====] - 10s 65ms/step - loss: 0.0190 - acc: 0.9951

```



In [138...

```
# Results
print ('Train Results:', results_train)
print ('Test Results:', results_test)
```

Train Results: [0.019027573987841606, 0.9950980544090271]
 Test Results: [0.1540239155292511, 0.9466292262077332]

In [139...

```
acc = history.history['acc']
val_acc = history.history['val_acc']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(len(acc))
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```



In [140...

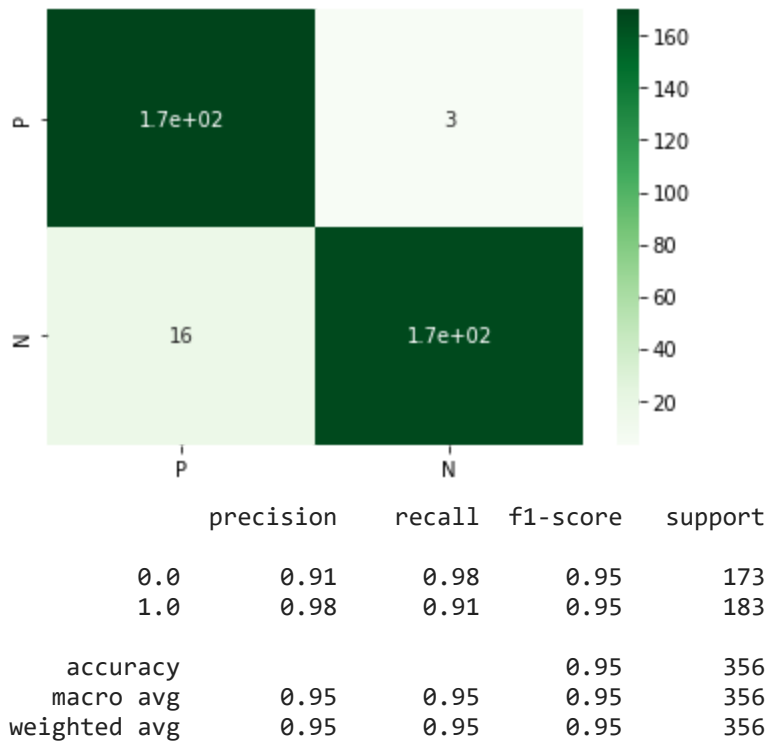
```
pred_y = (cnn.predict(test_images).ravel() > 0.5).astype(int)
```

In [141...

```
# Confusion Matrix
cm = confusion_matrix(test_y, pred_y)
f = sns.heatmap(cm, annot=True, cmap='Greens',
                xticklabels='PN', yticklabels='PN')
```

```
plt.show()


# Classification Report
print(classification_report(test_y, pred_y))
```



```
In [142... print('False Normal Rate:', (3/b2)*100)
```

False Normal Rate: 1.7341040462427744

```
In [143... plt.figure(1)
plt.plot([0, 1], [0, 1], 'k--')
plt.plot(fpr_, tpr_,
         label='model(area = {:.3f})'.format(auc_))
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc='best')
plt.show()
```




```
In [144... auc_
```

Out[144... 0.9909030607410215

```
In [145... pred_y = (cnn.predict(test_images).ravel() > auc_).astype(int)
```

```
In [146... # Confusion Matrix
cm = confusion_matrix(test_y, pred_y)
f = sns.heatmap(cm, annot=True, cmap='cubehelix',
                xticklabels='PN', yticklabels='PN')
plt.show()

# Classification Report
print(classification_report(test_y, pred_y))
```



	precision	recall	f1-score	support
0.0	0.74	0.99	0.85	173
1.0	0.99	0.67	0.80	183
accuracy			0.83	356
macro avg	0.87	0.83	0.82	356
weighted avg	0.87	0.83	0.82	356

```
In [147... print('False Normal Rate:', (1/b2)*100)
```

False Normal Rate: 0.5780346820809248

```
In [58]: cnn.save('XRAY_CNN_Model.h5')
```

Image Classification Process (CNN)

Importing Necessary Libraries

```
In [60]: from keras.models import load_model
from keras.preprocessing import image
from keras import models
import math
import numpy as np
import matplotlib.image as mpimg
import matplotlib.pyplot as plt
%matplotlib inline
```

Open/Load a Model

```
In [148... model = load_model('XRAY_CNN_Model.h5')
model.summary()
```

Model: "sequential_2"


Layer (type)	Output Shape	Param #
=====		
conv2d_3 (Conv2D)	(None, 148, 148, 32)	896

max_pooling2d_3 (MaxPooling2)	(None, 74, 74, 32)	0
conv2d_4 (Conv2D)	(None, 71, 71, 32)	16416
max_pooling2d_4 (MaxPooling2)	(None, 35, 35, 32)	0
conv2d_5 (Conv2D)	(None, 33, 33, 64)	18496
max_pooling2d_5 (MaxPooling2)	(None, 16, 16, 64)	0
flatten_1 (Flatten)	(None, 16384)	0
dense_6 (Dense)	(None, 64)	1048640
dense_7 (Dense)	(None, 1)	65
=====		
Total params: 1,084,513		
Trainable params: 1,084,513		
Non-trainable params: 0		

Open/Load Sample Test Image

In [149..

```
filename = 'data/test/normal/NORMAL-7725506-0001.jpeg'
img = image.load_img(filename, target_size=(150, 150))
plt.imshow(img)
plt.show()
```



View Image as Tensor

In [150..


```
img_tensor = image.img_to_array(img)
img_tensor = np.expand_dims(img_tensor, axis=0)

# Follow the Original Model Preprocessing
img_tensor /= 255.

# Check tensor shape
print(img_tensor.shape)

# Preview an image
plt.imshow(img_tensor[0])
plt.show()
```

(1, 150, 150, 3)



Visualization of Activation Layers

In [151...

```
# Extract model layer outputs
layer_outputs = [
    layer.output for layer in model.layers[:6]]

# Create a model for displaying the feature maps
activation_model = models.Model(
    inputs=model.input, outputs=layer_outputs)

activations = activation_model.predict(img_tensor)

# Extract Layer Names for Labelling
layer_names = []
for layer in model.layers[:6]:
    layer_names.append(layer.name)

total_features = sum([a.shape[-1] for a in activations])
total_features

n_cols = 12
n_rows = math.ceil(total_features / n_cols)

iteration = 0
fig, axes = plt.subplots(nrows=n_rows, ncols=n_cols,
                        figsize=(n_cols, n_rows*1.5))

for layer_n, layer_activation in enumerate(activations):
    n_channels = layer_activation.shape[-1]
    for ch_idx in range(n_channels):
        row = iteration // n_cols
        column = iteration % n_cols

        ax = axes[row, column]

        channel_image = layer_activation[0,
                                         :, :,
                                         ch_idx]

        channel_image -= channel_image.mean()
        channel_image /= channel_image.std()
        channel_image *= 32
        channel_image += 64
        channel_image = np.clip(
            channel_image, 0, 255).astype('uint8')

        ax.imshow(channel_image, aspect='auto',
                  cmap='viridis')
```

```

ax.get_xaxis().set_ticks([])
ax.get_yaxis().set_ticks([])


if ch_idx == 0:
    ax.set_title(layer_names[layer_n], fontsize=10)
    iteration += 1

fig.subplots_adjust(hspace=1.25)
plt.savefig('Intermediate_Activations_Visualized.pdf')
plt.show()



```

<ipython-input-151-2441043dfcd8>:40: RuntimeWarning: invalid value encountered in true_divide
channel_image /= channel_image.std()



conv2d_3





max_pooling2d_3




conv2d_4



max_pooling2d_4



conv2d_5



max_pooling2d_5

